

# Decision Trees on Amazon Food Reviews

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. index
2. Id
3. ProductId - unique identifier for the product
4. UserId - unique identifier for the user
5. ProfileName
6. HelpfulnessNumerator - number of users who found the review helpful
7. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
8. Score - rating between 1 and 5
9. Time - timestamp for the review
10. Summary - brief summary of the review
11. Text - text of the review
12. ProcessedText - Cleaned & Preprocessed Text of the review

**Objective: Given Amazon Food reviews, convert all the reviews into a vector using three techniques:**

- 1. Average W2V.**
- 2. Average TFIDF-W2V.**
- 3. GLoVe(Pre-trained).**

**Then perform following tasks under each technique:**

**Task 1. Split train and test data in a ratio of 80:20.**

**Task 2. Perform GridSearch Cross Validation to find optimal depth of decision tree.**

**Task 3. Apply DecisionTreeClassifier and report accuracy. Also check for train error.**

**Task 4. Plot decision tree using graphviz.**

[Q] How to determine if a review is positive or negative?

[Ans] We could use the Score/Rating. A rating of 4 or 5 could be considered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is neutral and ignored. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

Loading the data

SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently. Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [3]: import sqlite3
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from gensim.models import Word2Vec
import gensim
import csv
import re
import graphviz
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from sklearn.cross_validation import train_test_split
from sklearn.grid_search import GridSearchCV
```

```
In [3]: connection = sqlite3.connect('FinalAmazonFoodReviewsDataset.sqlite')
```

```
In [4]: data = pd.read_sql_query("SELECT * FROM Reviews", connection)
```

In [5]: data.head()

Out[5]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0										
	0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	Positive	1303862400	Good Quality Dog Food
1										
	1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	Negative	1346976000	Not as Advertised
2										
	2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	Positive	1219017600	"Delight" says it all
3										
	4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	Positive	1350777600	Great taffy
4										
	5	6	B006K2ZZ7K	ADT0SRK1MGOEU	Twoapennything	0	0	Positive	1342051200	Nice Taffy

In [6]: data.shape

Out[6]: (364171, 12)

```
In [7]: data["Score"].value_counts()
```

```
Out[7]: Positive    307061  
        Negative    57110  
        Name: Score, dtype: int64
```

```
In [8]: def changingScores(score):  
        if score == "Positive":  
            return 1  
        else:  
            return 0
```

```
In [9]: # changing score  
        # Positive = 1  
        # Negative = 0  
        actualScore = list(data["Score"])  
        positiveNegative = list(map(changingScores, actualScore)) #map(function, list of numbers)  
        data['Score'] = positiveNegative
```

In [10]: data.head()

Out[10]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0										
	0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality Dog Food
1										
	1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertised
2										
	2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all
3										
	4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	1	1350777600	Great taffy
4										
	5	6	B006K2ZZ7K	ADT0SRK1MG0EU	Twoapennything	0	0	1	1342051200	Nice Taffy

In [11]: allPositiveReviews = data[(data["Score"] == 1)]

```
In [12]: allPositiveReviews.shape
```

```
Out[12]: (307061, 12)
```

```
In [13]: positiveReviews_5000 = allPositiveReviews[:5000]
```

```
In [14]: positiveReviews_5000.shape
```

```
Out[14]: (5000, 12)
```

In [15]: `positiveReviews_5000.head()`

Out[15]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0										
	0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality Dog Food
2										
	2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all
3										
	4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	1	1350777600	Great taffy
4										
	5	6	B006K2ZZ7K	ADT0SRK1MG0EU	Twoapennything	0	0	1	1342051200	Nice Taffy
5										
	6	7	B006K2ZZ7K	A1SP2KVKFXXRU1	David C. Sullivan	0	0	1	1340150400	Great! Just as good as the expensive brands!

In [16]: `allNegativeReviews = data[(data["Score"] == 0)]`

```
In [17]: allNegativeReviews.shape
```

```
Out[17]: (57110, 12)
```

```
In [18]: negativeReviews_5000 = allNegativeReviews[:5000]
```

```
In [19]: negativeReviews_5000.shape
```

```
Out[19]: (5000, 12)
```



```
In [20]: negativeReviews_5000.head()
```

```
Out[20]:
```

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	T
1											
	1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertised	Prod arriv labe  Jun Sal Peanu
11											
	12	13	B0009XLVG0	A327PCT23YH90	LT	1	1	0	1339545600	My Cats Are Not Fans of the New Food	My c hi be hap eal Felic Pla
15											
	16	17	B001GVISJM	A3KLWF6WQ5BNYO	Erica Neathery	0	0	0	1348099200	poor taste	I l eat th and tl are gr
25											
	26	27	B001GVISJM	A3RXAU2N8KV45G	lady21	0	1	0	1332633600	Nasty No flavor	can just r No fla . J pla
45											
	47	51	B001EO5QW8	A108P30XVUFKXY	Roberto A	0	7	0	1203379200	Don't like it	T oatm is good. mus so

```
In [21]: frames_10000 = [positiveReviews_5000, negativeReviews_5000]
```

```
In [22]: FinalPositiveNegative = pd.concat(frames_10000)
```

```
In [23]: FinalPositiveNegative.shape
```

```
Out[23]: (10000, 12)
```

```
In [24]: #Sorting FinalDataframe by "Time"  
FinalSortedPositiveNegative_10000 = FinalPositiveNegative.sort_values('Time', axis=0, ascending=True, inplace=False)
```

```
In [25]: FinalSortedPositiveNegativeScore_10000 = FinalSortedPositiveNegative_10000["Score"]
```

```
In [26]: FinalSortedPositiveNegative_10000.shape
```

```
Out[26]: (10000, 12)
```

```
In [27]: FinalSortedPositiveNegativeScore_10000.shape
```

```
Out[27]: (10000,)
```

```
In [28]: Data = FinalSortedPositiveNegative_10000
```

```
In [29]: Data_Labels = FinalSortedPositiveNegativeScore_10000
```

```
In [30]: print(Data.shape)  
print(Data_Labels.shape)
```

```
(10000, 12)  
(10000,)
```

In [31]: Data.head()

Out[31]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
<b>772</b>										
	1146	1245	B00002Z754	A29Z5PI9BW2PU3	Robbie	7	7	1	961718400	Great Product
<b>771</b>										
	1145	1244	B00002Z754	A3B8RCEI0FXFI6	B G Chase	10	10	1	962236800	WOW Make your own 'slickers' !
<b>5822</b>										
	7427	8111	B0000EIE2Z	A3M174IC0VXOS2	Gail Cooke	3	3	1	1075420800	BEST BLUEBERRIES
<b>2418</b>										
	3481	3783	B00016UX0K	AF1PV3DIC0XM7	Robert Ashton	1	2	1	1081555200	Classic Condiment
<b>5206</b>										
	6790	7432	B0001E1IME	A2IKCTD1I73PLW	Adeba	2	8	1	1083456000	amazon monopoly/ripoff

## 1. Avg W2V

```
In [32]: i = 0
listOfSentences = []
for sentence in Data["ProcessedText"].values:
    subSentence = []
    for word in sentence.split():
        subSentence.append(word)

    listOfSentences.append(subSentence)
```

```
In [33]: print(Data['ProcessedText'].values[0])
print("\n")
print(listOfSentences[0:2])
print("\n")
print(type(listOfSentences))
```

this was realli good idea and the final product outstand use the decal car window and everybodi ask where bought the de  
cal made two thumb

```
[['this', 'was', 'realli', 'good', 'idea', 'and', 'the', 'final', 'product', 'outstand', 'use', 'the', 'decal', 'car',
'window', 'and', 'everybodi', 'ask', 'where', 'bought', 'the', 'decal', 'made', 'two', 'thumb'], ['just', 'receiv', 'sh
ipment', 'and', 'could', 'hard', 'wait', 'tri', 'this', 'product', 'love', 'which', 'what', 'call', 'them', 'instead',
'sticker', 'becaus', 'they', 'can', 'remov', 'easili', 'daughter', 'design', 'sign', 'print', 'revers', 'use', 'her',
'car', 'window', 'they', 'print', 'beauti', 'have', 'the', 'print', 'shop', 'program', 'go', 'have', 'lot', 'fun', 'wit
h', 'this', 'product', 'becaus', 'there', 'are', 'window', 'everywher', 'and', 'other', 'surfac', 'like', 'screen', 'an
d', 'comput', 'monitor']]
```

```
<class 'list'>
```

```
In [34]: w2vModel = gensim.models.Word2Vec(listOfSentences, size=300, min_count=5, workers=4)
```

In [35]: *# compute average word2vec for each review.*

```

sentenceAsW2V = []
for sentence in listOfSentences:
    sentenceVector = np.zeros(300)
    TotalWordsPerSentence = 0
    for word in sentence:
        try:
            vect = w2vModel.wv[word]
            sentenceVector += vect
            TotalWordsPerSentence += 1
        except:
            pass
    sentenceVector /= TotalWordsPerSentence
    sentenceAsW2V.append(sentenceVector)

print(type(sentenceAsW2V))
print(len(sentenceAsW2V))
print(len(sentenceAsW2V[0]))

```

```

<class 'list'>
10000
300

```

### Task 1. Split train and test data in a ratio of 80:20.

In [36]: `train_AvgW2V, test_AvgW2V, train_labels_AvgW2V, test_labels_AvgW2V = train_test_split(sentenceAsW2V, Data_Labels, test_si`

In [37]: `len(train_AvgW2V), len(test_AvgW2V), len(train_labels_AvgW2V), len(test_labels_AvgW2V)`

Out[37]: (8000, 2000, 8000, 2000)

### Task 2. Perform GridSearch Cross Validation to find optimal depth of decision tree.

```
In [38]: clf = DecisionTreeClassifier(min_samples_split = 4, min_samples_leaf = 4)

hyper_parameters = {'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 11, 13, 15, 17]}
bestCV = GridSearchCV(clf, hyper_parameters, scoring = "accuracy", cv = 3)
bestCV.fit(train_AvgW2V, train_labels_AvgW2V)

print(bestCV.best_estimator_)
```

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=5,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=4, min_samples_split=4,
                        min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                        splitter='best')
```

```
In [39]: best_parameter = bestCV.best_params_
best_parameter["max_depth"]
```

```
Out[39]: 5
```

```
In [40]: scoreData = bestCV.grid_scores_
scoreData
```

```
Out[40]: [mean: 0.70037, std: 0.01147, params: {'max_depth': 1},
mean: 0.70950, std: 0.00853, params: {'max_depth': 2},
mean: 0.72288, std: 0.01223, params: {'max_depth': 3},
mean: 0.73275, std: 0.00629, params: {'max_depth': 4},
mean: 0.73550, std: 0.00935, params: {'max_depth': 5},
mean: 0.73088, std: 0.00579, params: {'max_depth': 6},
mean: 0.71837, std: 0.00401, params: {'max_depth': 7},
mean: 0.70862, std: 0.00808, params: {'max_depth': 8},
mean: 0.70937, std: 0.00697, params: {'max_depth': 9},
mean: 0.69737, std: 0.00369, params: {'max_depth': 11},
mean: 0.68937, std: 0.00570, params: {'max_depth': 13},
mean: 0.69475, std: 0.00222, params: {'max_depth': 15},
mean: 0.69025, std: 0.00659, params: {'max_depth': 17}]
```

```

In [41]: error = []
parameter = []
for i in range(len(scoreData)):
    error.append(1 - scoreData[i][1])
    parameter.append(scoreData[i][0]["max_depth"])

plt.plot(parameter, np.round(error, 4))
plt.ylim(ymin=0.2)
plt.ylim(ymax=0.4)
plt.xlabel("Maximum Depth", fontsize=16)
plt.ylabel("Error", fontsize=16)

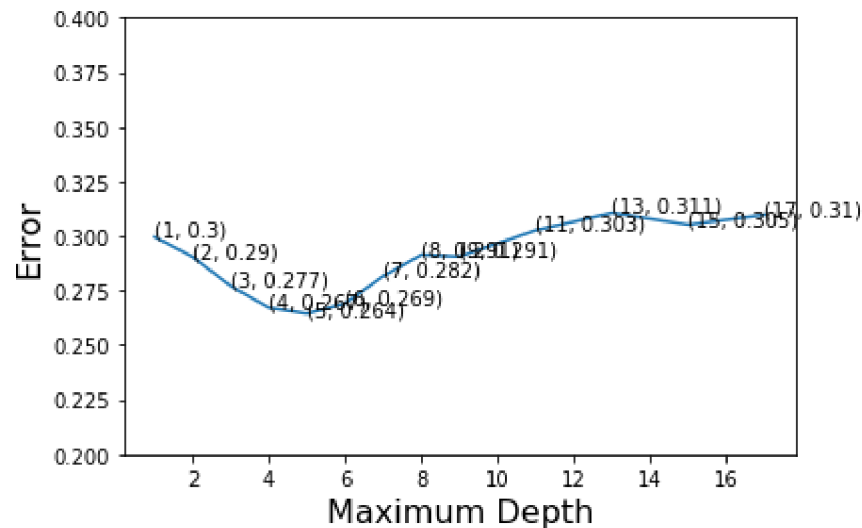
error1 = []
for e in error:
    error1.append(np.round(e,3))

parameter1 = []
for p in parameter:
    parameter1.append(np.round(p,3))

for xy in zip(parameter1, error1):
    plt.annotate(xy,xy)

plt.show()

```



### Task 3. Apply DecisionTreeClassifier and report accuracy. Also check for train error.

```
In [42]: clf_AvgW2V = DecisionTreeClassifier(max_depth = best_parameter["max_depth"], min_samples_split = 4, min_samples_leaf = 4)

clf_AvgW2V = clf_AvgW2V.fit(train_AvgW2V, train_labels_AvgW2V)

prediction_AvgW2V = clf_AvgW2V.predict(test_AvgW2V)

AccuracyScore_AvgW2V = accuracy_score(test_labels_AvgW2V, prediction_AvgW2V) * 100

print("Accuracy score of decision tree = "+str(AccuracyScore_AvgW2V)+"%")

Accuracy score of decision tree = 73.2%
```

```
In [43]: model_AvgW2V_tr = DecisionTreeClassifier(max_depth = best_parameter["max_depth"], min_samples_split = 4, min_samples_leaf = 4)

model_AvgW2V_tr.fit(train_AvgW2V, train_labels_AvgW2V)

prediction_AvgW2V_tr = model_AvgW2V_tr.predict(train_AvgW2V)

AccuracyScore_AvgW2V_tr = accuracy_score(train_labels_AvgW2V, prediction_AvgW2V_tr)

print("Train error of decision tree = "+str(1- AccuracyScore_AvgW2V_tr))

Train error of decision tree = 0.21987500000000004
```

```
In [44]: Confusion_Matrix = confusion_matrix(test_labels_AvgW2V, prediction_AvgW2V)

print("Confusion Matrix on L2 regularization \n"+str(Confusion_Matrix))

Confusion Matrix on L2 regularization
[[727 250]
 [286 737]]
```

```
In [45]: tn, fp, fn, tp = confusion_matrix(test_labels_AvgW2V, prediction_AvgW2V).ravel()

tn, fp, fn, tp
```

Out[45]: (727, 250, 286, 737)



## Task 4. Plot decision tree using graphviz.

```
In [65]: dot_data = tree.export_graphviz(clf_AvgW2V, out_file=None, filled=True, rounded=True, special_characters=True)
graph = graphviz.Source(dot_data)
graph.render("AvgW2V_DecisionTree_Graph")
```

```
Out[65]: 'AvgW2V_DecisionTree_Graph.pdf'
```

## 2. Average TFIDF-W2V.

```
In [49]: tfidf_vect = TfidfVectorizer(ngram_range = (1,2))
tfidf = tfidf_vect.fit_transform(Data["ProcessedText"].values)
```

```
In [50]: w2v_Model = gensim.models.Word2Vec(listOfSentences, size=300, min_count=5, workers=4)
```

```
In [51]: print(tfidf.shape)
print(type(tfidf))

(10000, 237703)
<class 'scipy.sparse.csr.csr_matrix'>
```

```

In [52]: # TF-IDF weighted Word2Vec
tfidf_features = tfidf_vect.get_feature_names()

tfidf_w2v = []
reviews = 0

for sentence in listOfSentences:
    sentenceVector = np.zeros(300)
    weightTfidfSum = 0
    for word in sentence:
        try:
            W2V_Vector = w2v_Model.wv[word]
            tfidfVector = tfidf[reviews, tfidf_features.index(word)]
            sentenceVector += (W2V_Vector * tfidfVector)
            weightTfidfSum += tfidfVector
        except:
            pass
    sentenceVector /= weightTfidfSum
    tfidf_w2v.append(sentenceVector)
    reviews += 1

```

### Task 1. Split train and test data in a ratio of 80:20

```

In [53]: train_tfidf_w2v, test_tfidf_w2v, train_labels_tfidf_w2v, test_labels_tfidf_w2v = train_test_split(tfidf_w2v, Data_Labels,

```

```

In [54]: len(train_tfidf_w2v), len(test_tfidf_w2v), len(train_labels_tfidf_w2v), len(test_labels_tfidf_w2v)

```

```

Out[54]: (8000, 2000, 8000, 2000)

```

### Task 2. Perform GridSearch Cross Validation to find optimal depth of decision tree.

```
In [56]: clf = DecisionTreeClassifier(min_samples_split = 4, min_samples_leaf = 4)

hyper_parameters = {'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 11, 13, 15, 17]}
bestCV = GridSearchCV(clf, hyper_parameters, scoring = "accuracy", cv = 3)
bestCV.fit(train_tfidf_w2v, train_labels_tfidf_w2v)

print(bestCV.best_estimator_)
```

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=5,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=4, min_samples_split=4,
                        min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                        splitter='best')
```

```
In [57]: best_parameter = bestCV.best_params_
best_parameter["max_depth"]
```

```
Out[57]: 5
```

```
In [58]: scoreData = bestCV.grid_scores_
scoreData
```

```
Out[58]: [mean: 0.67325, std: 0.00715, params: {'max_depth': 1},
mean: 0.67550, std: 0.00439, params: {'max_depth': 2},
mean: 0.68537, std: 0.00670, params: {'max_depth': 3},
mean: 0.69650, std: 0.00656, params: {'max_depth': 4},
mean: 0.70037, std: 0.00464, params: {'max_depth': 5},
mean: 0.69763, std: 0.00508, params: {'max_depth': 6},
mean: 0.69675, std: 0.00997, params: {'max_depth': 7},
mean: 0.68137, std: 0.00942, params: {'max_depth': 8},
mean: 0.67312, std: 0.01638, params: {'max_depth': 9},
mean: 0.67050, std: 0.01209, params: {'max_depth': 11},
mean: 0.66825, std: 0.01034, params: {'max_depth': 13},
mean: 0.66212, std: 0.01194, params: {'max_depth': 15},
mean: 0.66325, std: 0.00999, params: {'max_depth': 17}]
```

```

In [59]: error = []
parameter = []
for i in range(len(scoreData)):
    error.append(1 - scoreData[i][1])
    parameter.append(scoreData[i][0]["max_depth"])

plt.plot(parameter, np.round(error, 4))
plt.ylim(ymin=0.2)
plt.ylim(ymax=0.4)
plt.xlabel("Maximum Depth", fontsize=16)
plt.ylabel("Error", fontsize=16)


error1 = []
for e in error:
    error1.append(np.round(e,3))

parameter1 = []
for p in parameter:
    parameter1.append(np.round(p,3))

for xy in zip(parameter1, error1):
    plt.annotate(xy,xy)

plt.show()

```



### Task 3. Apply DecisionTreeClassifier and report accuracy. Also check for train error

```
In [60]: clf_TFIDF_W2V = DecisionTreeClassifier(max_depth = best_parameter["max_depth"], min_samples_split = 4, min_samples_leaf = 1)
clf_TFIDF_W2V = clf_TFIDF_W2V.fit(train_tfidf_w2v, train_labels_tfidf_w2v)
prediction_TFIDF_W2V = clf_TFIDF_W2V.predict(test_tfidf_w2v)
AccuracyScore_TFIDF_W2V = accuracy_score(test_labels_tfidf_w2v, prediction_TFIDF_W2V) * 100
print("Accuracy score of decision tree = "+str(AccuracyScore_TFIDF_W2V)+"%")
```

Accuracy score of decision tree = 70.85000000000001%

```
In [96]: clf_TFIDF_W2V_tr = DecisionTreeClassifier(max_depth = best_parameter["max_depth"], min_samples_split = 4, min_samples_leaf = 1)
clf_TFIDF_W2V_tr = clf_TFIDF_W2V_tr.fit(train_tfidf_w2v, train_labels_tfidf_w2v)
prediction_TFIDF_W2V_tr = clf_TFIDF_W2V_tr.predict(train_tfidf_w2v)
AccuracyScore_TFIDF_W2V_tr = accuracy_score(train_labels_tfidf_w2v, prediction_TFIDF_W2V_tr)
print("Train error of decision tree = "+str(1- AccuracyScore_TFIDF_W2V_tr))
```

Train error of decision tree = 0.25712500000000005

```
In [62]: Confusion_Matrix = confusion_matrix(test_labels_tfidf_w2v, prediction_TFIDF_W2V)
print("Confusion Matrix on L2 regularization \n"+str(Confusion_Matrix))
```

Confusion Matrix on L2 regularization

```
[[717 300]
 [283 700]]
```

```
In [63]: tn, fp, fn, tp = confusion_matrix(test_labels_AvgW2V, prediction_AvgW2V).ravel()
tn, fp, fn, tp
```

Out[63]: (727, 250, 286, 737)

## Task 4. Plot decision tree using graphviz.

```
In [64]: dot_data = tree.export_graphviz(clf_TFIDF_W2V, out_file=None, filled=True, rounded=True, special_characters=True)
graph = graphviz.Source(dot_data)
graph.render("TFIDF_W2V_DecisionTree_Graph")
```

```
Out[64]: 'TFIDF_W2V_DecisionTree_Graph.pdf'
```

## 3. GLoVe

```
In [32]: def cleanhtml(sentence): #function to clean htmltags
        cleanr = re.compile("<.*?>")
        cleantext = re.sub(cleanr, " ", sentence)
        return cleantext

        def cleanpunc(sentence): #function to clean the word of any punctuation or special characters
        cleaned = re.sub(r'[?|!|\\'|"|#]', '', sentence)
        cleaned = re.sub(r'[,|,|)|(|\\|/]', ' ', cleaned)
        return cleaned
```

```
In [33]: #removing HTML tags and punctuation from our text

i = 0
final_string = []
s = ""
for sentence in data["Text"].values:
    filteredSentence = []
    EachReviewText = ""
    sentenceHTMLCleaned = cleanhtml(sentence)
    for eachWord in sentenceHTMLCleaned.split():
        for sentencePunctCleaned in cleanpunc(eachWord).split():
            if((sentencePunctCleaned.isalpha()) & (len(sentencePunctCleaned)>2)):
                sentenceLower = sentencePunctCleaned.lower()
                filteredSentence.append(sentenceLower)

    EachReviewText = ' '.join(filteredSentence)
    final_string.append(EachReviewText)
```

```
In [34]: data["ProcessedText2"] = final_string
```

```
In [35]: data.head()
```

```
Out[35]:
```

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0										
	0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality Dog Food
1										
	1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertised
2										
	2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all
3										
	4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	1	1350777600	Great taffy
4										
	5	6	B006K2ZZ7K	ADT0SRK1MG0EU	Twoapennything	0	0	1	1342051200	Nice Taffy

```
In [36]: allPositiveReviews2 = data[(data["Score"] == 1)]
```

```
In [37]: allPositiveReviews2.shape
```

```
Out[37]: (307061, 13)
```

```
In [38]: positiveReviews2_500 = allPositiveReviews2[:500]
```

```
In [39]: positiveReviews2_500.shape
```

```
Out[39]: (500, 13)
```



```
In [40]: positiveReviews2_500.head()
```

```
Out[40]:
```

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0										
	0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality Dog Food
2										
	2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all
3										
	4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	1	1350777600	Great taffy
4										
	5	6	B006K2ZZ7K	ADT0SRK1MG0EU	Twoapennything	0	0	1	1342051200	Nice Taffy
5										
	6	7	B006K2ZZ7K	A1SP2KVKFXXRU1	David C. Sullivan	0	0	1	1340150400	Great! Just as good as the expensive brands!

```
In [41]: allNegativeReviews2 = data[(data["Score"] == 0)]
```

```
In [42]: allNegativeReviews2.shape
```

```
Out[42]: (57110, 13)
```

```
In [43]: negativeReviews2_500 = allNegativeReviews2[:500]
```

```
In [44]: negativeReviews2_500.shape
```

```
Out[44]: (500, 13)
```

In [45]: `negativeReviews2_500.head()`

Out[45]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	T
1											Prod arriv labe
	1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertised	Jun Sal Peanu
11											My c hi be hap eal Felic Pla
	12	13	B0009XLVG0	A327PCT23YH90	LT	1	1	0	1339545600	My Cats Are Not Fans of the New Food	
15											I l eat th and tl are gr
	16	17	B001GVISJM	A3KLWF6WQ5BNYO	Erica Neathery	0	0	0	1348099200	poor taste	
25											watc -
	26	27	B001GVISJM	A3RXAU2N8KV45G	lady21	0	1	0	1332633600	Nasty No flavor	cand just re No fla . J pla
45											T oatm is good. mus so
	47	51	B001EO5QW8	A108P30XVUFKXY	Roberto A	0	7	0	1203379200	Don't like it	

In [46]: `frames2_1000 = [positiveReviews2_500, negativeReviews2_500]`

```
In [47]: FinalPositiveNegative2 = pd.concat(frames2_1000)
```

```
In [48]: FinalPositiveNegative2.shape
```

```
Out[48]: (1000, 13)
```

```
In [49]: #Sorting FinalDataframe by "Time"  
FinalSortedPositiveNegative2_1000 = FinalPositiveNegative2.sort_values('Time', axis=0, ascending=True, inplace=False)
```

```
In [50]: FinalSortedPositiveNegativeScore2_1000 = FinalSortedPositiveNegative2_1000["Score"]
```

```
In [51]: FinalSortedPositiveNegative2_1000.shape
```

```
Out[51]: (1000, 13)
```

```
In [52]: FinalSortedPositiveNegativeScore2_1000.shape
```

```
Out[52]: (1000,)
```

```
In [53]: Data2 = FinalSortedPositiveNegative2_1000
```

```
In [54]: Data2_Labels = FinalSortedPositiveNegativeScore2_1000
```

```
In [55]: print(Data2.shape)  
print(Data2_Labels.shape)
```

```
(1000, 13)  
(1000,)
```

In [56]: Data2.head()

Out[56]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	
9	10	11	B0001PB9FE	A3HDKO7OW0QNK4	Canadian Fan	1	1	1	1107820800	Th S
1653	2106	2296	B0001VWE02	AQM74O8Z4FMS0	Sunshine	0	0	0	1127606400	Belov
2558	3667	3984	B0005ZHWXI	A26HFSVLAGULIM	Heather L. Parisi "Robert and Heather Parisi"	0	1	0	1131235200	BLUe
1341	1779	1935	B000F4EU52	A2PNOU7NXB1JE4	Peggy "pab920"	14	17	0	1153008000	In
2307	3362	3661	B000FDKQC4	A1PNP10DP0M7V1	D. Chamberlain "dchamberlain072002"	7	8	0	1156377600	.

```
In [57]: i = 0
listOfSentences2 = []
for sentence in Data2["ProcessedText2"].values:
    subSentence = []
    for word in sentence.split():
        subSentence.append(word)

    listOfSentences2.append(subSentence)
```

```
In [58]: print(Data2['ProcessedText2'].values[0])
print("\n")
print(listOfSentences2[0:2])
print("\n")
print(type(listOfSentences2))
print("\n")
print(len(listOfSentences2))
```

dont know its the cactus the tequila just the unique combination ingredients but the flavour this hot sauce makes one kind picked bottle once trip were and brought back home with and were totally blown away when realized that simply couldnt find anywhere our city were bummed now because the magic the internet have case the sauce and are ecstatic because you love hot sauce mean really love hot sauce but dont want sauce that tastelessly burns your throat grab bottle tequila picante gourmet inclan just realize that once you taste you will never want use any other sauce thank you for the personal incredible service

```
[['dont', 'know', 'its', 'the', 'cactus', 'the', 'tequila', 'just', 'the', 'unique', 'combination', 'ingredients', 'but', 'the', 'flavour', 'this', 'hot', 'sauce', 'makes', 'one', 'kind', 'picked', 'bottle', 'once', 'trip', 'were', 'and', 'brought', 'back', 'home', 'with', 'and', 'were', 'totally', 'blown', 'away', 'when', 'realized', 'that', 'simply', 'couldnt', 'find', 'anywhere', 'our', 'city', 'were', 'bummed', 'now', 'because', 'the', 'magic', 'the', 'internet', 'have', 'case', 'the', 'sauce', 'and', 'are', 'ecstatic', 'because', 'you', 'love', 'hot', 'sauce', 'mean', 'really', 'love', 'hot', 'sauce', 'but', 'dont', 'want', 'sauce', 'that', 'tastelessly', 'burns', 'your', 'throat', 'grab', 'bottle', 'tequila', 'picante', 'gourmet', 'inclan', 'just', 'realize', 'that', 'once', 'you', 'taste', 'you', 'will', 'never', 'want', 'use', 'any', 'other', 'sauce', 'thank', 'you', 'for', 'the', 'personal', 'incredible', 'service'], ['too', 'much', 'the', 'white', 'pith', 'this', 'orange', 'peel', 'making', 'the', 'product', 'overly', 'bitter', 'and', 'diluting', 'the', 'real', 'good', 'taste', 'the', 'orange', 'zest']]
```

```
<class 'list'>
```

```
1000
```

```
In [59]: #Loading pre-trained GloVe vectors
words = pd.read_table("glove.6B.100d.txt", sep=" ", index_col=0, header=None, quoting=csv.QUOTE_NONE)

# Here, We have downloaded pre-trained Glove vectors. You just have to type "Glove word vectors" on google then click on
# "https://nlp.stanford.edu/projects/glove/" Link. Then you can download pre-trained word-vectors. Zip file will be
# downloaded, you just have to extract it then load the txt file from extracted folder into ipython notebook using pandas
# just like we have done above.
```

```
In [60]: def check(word):
         if (words.index == word).any():
             return 1
         else:
             return 0
```

```
In [61]: # compute average GloVe for each review.
sentenceAsGlove = []
for sentence in listOfSentences2:
    sentenceVector = np.zeros(100)
    TotalWordsPerSentence = 0
    for word in sentence:
        if check(word) == 1:
            vect = words.loc[word]
            sentenceVector += vect
            TotalWordsPerSentence += 1

    sentenceVector /= TotalWordsPerSentence
    sentenceAsGlove.append(sentenceVector)

print(type(sentenceAsGlove))
print(len(sentenceAsGlove))
print(len(sentenceAsGlove[0]))

<class 'list'>
1000
100
```

### Task 1. Split train and test data in a ratio of 80:20

```
In [62]: train_GLoVe, test_GLoVe, train_labels_GLoVe, test_labels_GLoVe = train_test_split(sentenceAsGlove, Data2_Labels, test_size=0.2)
```

```
In [63]: len(train_GLoVe), len(test_GLoVe), len(train_labels_GLoVe), len(test_labels_GLoVe)
```

```
Out[63]: (800, 200, 800, 200)
```

### Task 2. Perform GridSearch Cross Validation to find optimal depth of decision tree.

```
In [83]: clf = DecisionTreeClassifier(min_samples_split = 4, min_samples_leaf = 4)

hyper_parameters = {'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 11, 13, 15, 17]}
bestCV = GridSearchCV(clf, hyper_parameters, scoring = "accuracy", cv = 3)
bestCV.fit(train_GLoVe, train_labels_GLoVe)

print(bestCV.best_estimator_)
```

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=1,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=4, min_samples_split=4,
                        min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                        splitter='best')
```

```
In [84]: best_parameter = bestCV.best_params_
best_parameter["max_depth"]
```

```
Out[84]: 1
```

```
In [85]: scoreData = bestCV.grid_scores_
scoreData
```

```
Out[85]: [mean: 0.65625, std: 0.00854, params: {'max_depth': 1},
mean: 0.63125, std: 0.01595, params: {'max_depth': 2},
mean: 0.65250, std: 0.03836, params: {'max_depth': 3},
mean: 0.63125, std: 0.02002, params: {'max_depth': 4},
mean: 0.62250, std: 0.01838, params: {'max_depth': 5},
mean: 0.62875, std: 0.03259, params: {'max_depth': 6},
mean: 0.62875, std: 0.02142, params: {'max_depth': 7},
mean: 0.63125, std: 0.03681, params: {'max_depth': 8},
mean: 0.63750, std: 0.01427, params: {'max_depth': 9},
mean: 0.62875, std: 0.01893, params: {'max_depth': 11},
mean: 0.62625, std: 0.01323, params: {'max_depth': 13},
mean: 0.63000, std: 0.02929, params: {'max_depth': 15},
mean: 0.62875, std: 0.02118, params: {'max_depth': 17}]
```



```

In [86]: error = []
parameter = []
for i in range(len(scoreData)):
    error.append(1 - scoreData[i][1])
    parameter.append(scoreData[i][0]["max_depth"])

plt.plot(parameter, np.round(error, 4))
plt.ylim(ymin=0.2)
plt.ylim(ymax=0.4)
plt.xlabel("Maximum Depth", fontsize=16)
plt.ylabel("Error", fontsize=16)


error1 = []
for e in error:
    error1.append(np.round(e,3))

parameter1 = []
for p in parameter:
    parameter1.append(np.round(p,3))

for xy in zip(parameter1, error1):
    plt.annotate(xy,xy)

plt.show()

```



### Task 3. Apply DecisionTreeClassifier and report accuracy. Also check for train error

```
In [92]: clf_GLoVe = DecisionTreeClassifier(max_depth = 5, min_samples_split = 4, min_samples_leaf = 4)

clf_GLoVe = clf_GLoVe.fit(train_GLoVe, train_labels_GLoVe)

prediction_GLoVe = clf_GLoVe.predict(test_GLoVe)

AccuracyScore_GLoVe = accuracy_score(test_labels_GLoVe, prediction_GLoVe) * 100

print("Accuracy score of decision tree = "+str(AccuracyScore_GLoVe)+"%")

Accuracy score of decision tree = 64.0%
```

```
In [93]: clf_GLoVe_tr = DecisionTreeClassifier(max_depth = 5, min_samples_split = 4, min_samples_leaf = 4)

clf_GLoVe_tr = clf_GLoVe_tr.fit(train_GLoVe, train_labels_GLoVe)

prediction_GLoVe_tr = clf_GLoVe_tr.predict(train_GLoVe)

AccuracyScore_GLoVe_tr = accuracy_score(train_labels_GLoVe, prediction_GLoVe_tr)

print("Train error of decision tree = "+str(1 - AccuracyScore_GLoVe_tr))

Train error of decision tree = 0.17000000000000004
```

```
In [94]: Confusion_Matrix = confusion_matrix(test_labels_GLoVe, prediction_GLoVe)

print("Confusion Matrix on L2 regularization \n"+str(Confusion_Matrix))

Confusion Matrix on L2 regularization
[[56 33]
 [39 72]]
```

```
In [95]: tn, fp, fn, tp = confusion_matrix(test_labels_GLoVe, prediction_GLoVe).ravel()
tn, fp, fn, tp
```

```
Out[95]: (56, 33, 39, 72)
```

## Task 4. Plot decision tree using graphviz.

```
In [96]: dot_data = tree.export_graphviz(clf_GLoVe, out_file=None, filled=True, rounded=True, special_characters=True)
graph = graphviz.Source(dot_data)
graph.render("GLoVe_DecisionTree_Graph")
```

```
Out[96]: 'GLoVe_DecisionTree_Graph.pdf'
```

## Summary

### Avg W2V

1. Optimal Value of depth from Grid Search = 5
2. Accuracy = 73.2%
3. Train Error = 0.2198

### TFIDF-W2V

1. Optimal Value of depth from Grid Search = 5
2. Accuracy = 70.85%
3. Train Error = 0.2571

### GLoVe

1. Optimal Value of depth from Grid Search = 1
2. Final considered value of depth = 5
3. Accuracy = 64%
4. Train Error = 0.170