

Naive Bayes on Amazon Food Reviews Dataset

Using Binary Bag of Words(BoW) and TFIDF Techniques

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. index
2. Id
3. ProductId - unique identifier for the product
4. UserId - unique identifier for the user
5. ProfileName
6. HelpfulnessNumerator - number of users who found the review helpful
7. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
8. Score - rating between 1 and 5
9. Time - timestamp for the review
10. Summary - brief summary of the review
11. Text - text of the review
12. ProcessedText - Cleaned & Preprocessed Text of the review

Objective: Given Amazon Food reviews, convert all the reviews into a vector using binary BoW and tfidf techniques and then apply forward chaining cross validation to determine the right value of alpha. Finally apply Naive Bayes on the top of it. Find top 10 important features. Also calculate followings performance metric scores: Accuracy, precision, recall, F1 Score and confusion metric.

[Q] How to determine if a review is positive or negative?

[Ans] We could use the Score/Rating. A rating of 4 or 5 could be considered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is neutral and ignored. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

Loading the data

SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently. Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: import sqlite3
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

from sklearn.naive_bayes import BernoulliNB, MultinomialNB

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix

from wordcloud import WordCloud

In [2]: connection = sqlite3.connect('FinalAmazonFoodReviewsDataset.sqlite')

In [3]: data = pd.read_sql_query("SELECT * FROM Reviews", connection)
```

```
In [4]: data.head()
```

```
Out[4]:
```

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0	0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	Positive	1303862400	Good Quality Dog Food
1	1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	Negative	1346976000	Not as Advertised
2	2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	Positive	1219017600	"Delight" says it all
3	4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	Positive	1350777600	Great taffy
4	5	6	B006K2ZZ7K	ADT0SRK1MG0EU	Twoapennything	0	0	Positive	1342051200	Nice Taffy



```
In [5]: data.shape
```

```
Out[5]: (364171, 12)
```

```
In [6]: data["Score"].value_counts()
```

```
Out[6]: Positive    307061  
        Negative    57110  
        Name: Score, dtype: int64
```

```
In [7]: def changingScores(score):  
        if score == "Positive":  
            return 1  
        else:  
            return 0
```

```
In [8]: # changing score  
        # Positive = 1  
        # Negative = 0  
        actualScore = list(data["Score"])  
        positiveNegative = list(map(changingScores, actualScore)) #map(function, list of numbers)  
        data['Score'] = positiveNegative
```

In [9]: data.head()

Out[9]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0	0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality Dog Food
1	1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertised
2	2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all
3	4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	1	1350777600	Great taffy
4	5	6	B006K2ZZ7K	ADT0SRK1MGOEU	Twoapennything	0	0	1	1342051200	Nice Taffy



In [10]: allPositiveReviews = data[(data["Score"] == 1)]

```
In [11]: allPositiveReviews.shape
```

```
Out[11]: (307061, 12)
```

```
In [12]: positiveReviews_30000 = allPositiveReviews[:30000]
```

```
In [13]: positiveReviews_30000.shape
```

```
Out[13]: (30000, 12)
```

In [14]: `positiveReviews_30000.head()`

Out[14]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0	0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality Dog Food
2	2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all
3	4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	1	1350777600	Great taffy
4	5	6	B006K2ZZ7K	ADT0SRK1MGOEU	Twoapennything	0	0	1	1342051200	Nice Taffy
5	6	7	B006K2ZZ7K	A1SP2KVKFXXRU1	David C. Sullivan	0	0	1	1340150400	Great! Just as good as the expensive brands!



In [15]: `allNegativeReviews = data[(data["Score"] == 0)]`

```
In [16]: allNegativeReviews.shape
```

```
Out[16]: (57110, 12)
```

```
In [17]: negativeReviews_30000 = allNegativeReviews[:30000]
```

```
In [18]: negativeReviews_30000.shape
```

```
Out[18]: (30000, 12)
```



```
In [19]: negativeReviews_30000.head()
```

Out[19]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	T	
	1	1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertised	Prod arriv labe
	11	12	13	B0009XLVG0	A327PCT23YH90	LT	1	1	0	1339545600	My Cats Are Not Fans of the New Food	Jun Sal Peanu My c hi be hap eal Felic Pla
	15	16	17	B001GVISJM	A3KLWF6WQ5BNYO	Erica Neathery	0	0	0	1348099200	poor taste	I k eal th and tl are ge
	25	26	27	B001GVISJM	A3RXAU2N8KV45G	lady21	0	1	0	1332633600	Nasty No flavor	watc - cand just re No fla . J pla
	45	47	51	B001EO5QW8	A108P30XVUFKXY	Roberto A	0	7	0	1203379200	Don't like it	T oatm is good. mus so

```
In [20]: frames_60000 = [positiveReviews_30000, negativeReviews_30000]
```

```
In [21]: FinalPositiveNegative = pd.concat(frames_60000)
```

```
In [22]: FinalPositiveNegative.shape
```

```
Out[22]: (60000, 12)
```

```
In [23]: #Sorting FinalDataframe by "Time"  
FinalSortedPositiveNegative_60000 = FinalPositiveNegative.sort_values('Time', axis=0, ascending=True, inplace=False)
```

```
In [24]: FinalSortedPositiveNegativeScore_60000 = FinalSortedPositiveNegative_60000["Score"]
```

```
In [25]: FinalSortedPositiveNegative_60000.shape
```

```
Out[25]: (60000, 12)
```

```
In [26]: FinalSortedPositiveNegativeScore_60000.shape
```

```
Out[26]: (60000,)
```

In [27]: FinalSortedPositiveNegative_60000.head()

Out[27]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summa
772	1146	1245	B00002Z754	A29Z5PI9BW2PU3	Robbie	7	7	1	961718400	Gre Produ
771	1145	1244	B00002Z754	A3B8RCEI0FXFI6	B G Chase	10	10	1	962236800	WOW Ma your ov 'slickers
19460	28087	30630	B00008RCMI	A284C7M23F0APC	A. Mendoza	0	0	1	1067040000	Be sugarle gum eve
19459	28086	30629	B00008RCMI	A19E94CF5O1LY7	Andrew Arnold	0	0	1	1067040000	I've chew this gu many time but user
85400	121056	131233	B00004RAMX	A1PYZPS1QYR036	Kazantzakis "hinterlands"	5	8	0	1067385600	Woodstrea Gopher Tr: 06



In [28]: Final_Data = FinalSortedPositiveNegative_60000

```
In [29]: Final_Data.shape
```

```
Out[29]: (60000, 12)
```

```
In [30]: Final_Data_Labels = FinalSortedPositiveNegativeScore_60000
```

```
In [31]: Final_Data_Labels.shape
```

```
Out[31]: (60000,)
```

Binary BoW

```
In [32]: positive_reviews = Final_Data[(Final_Data["Score"] == 1)]  
negative_reviews = Final_Data[(Final_Data["Score"] == 0)]
```

```
In [33]: positive_reviews.shape, negative_reviews.shape
```

```
Out[33]: ((30000, 12), (30000, 12))
```

```
In [34]: positive_bow_vect = CountVectorizer(stop_words = "english")  
positive_bow = positive_bow_vect.fit_transform(positive_reviews["ProcessedText"].values)
```

```
In [35]: positive_bow.shape
```

```
Out[35]: (30000, 20285)
```

```
In [36]: features_positive = positive_bow_vect.get_feature_names()  
len(features_positive), type(features_positive)
```

```
Out[36]: (20285, list)
```

```
In [37]: count = []  
for i in range(len(features_positive)):  
    total = positive_bow.getcol(i).sum() # it will give sum of all the values in 'i'th column  
    count.append(total)
```

```
In [38]: positive_dict = dict(zip(features_positive, count))
```

```
In [39]: sortedDict_Positive = sorted(positive_dict.items(), key = lambda positive_dict: positive_dict[1], reverse = True)
```

```
In [40]: for i in range(200):  
         print(sortedDict_Positive[i])
```

```
('like', 13895)  
('tast', 12635)  
('good', 11457)  
('love', 10510)  
('flavor', 10373)  
('great', 10208)  
('use', 9922)  
('veri', 8959)  
('tri', 8710)  
('product', 8691)  
('just', 8638)  
('tea', 8437)  
('coffe', 7759)  
('make', 7602)  
('food', 6632)  
('time', 5416)  
('eat', 5313)  
('buy', 5269)  
('realli', 5183)  
('onli', 5119)  
('dog', 5041)  
('price', 4705)  
('amazon', 4644)  
('best', 4612)  
('littl', 4501)  
('chocol', 4492)  
('dont', 4443)  
('drink', 4361)  
('order', 4336)  
('becaus', 4037)  
('ive', 3990)  
('mix', 3987)  
('treat', 3931)  
('store', 3818)  
('bag', 3785)  
('better', 3560)  
('recommend', 3489)  
('ani', 3423)
```

('day', 3333)
('year', 3306)
('sugar', 3303)
('sweet', 3164)
('cup', 3163)
('high', 3098)
('want', 2943)
('box', 2922)
('look', 2889)
('water', 2876)
('enjoy', 2875)
('work', 2838)
('brand', 2808)
('delici', 2797)
('bar', 2674)
('nice', 2664)
('favorit', 2618)
('cat', 2617)
('bit', 2607)
('think', 2605)
('way', 2603)
('add', 2590)
('need', 2583)
('packag', 2571)
('purchas', 2560)
('sinc', 2468)
('perfect', 2429)
('thing', 2389)
('bought', 2381)
('pack', 2348)
('lot', 2325)
('free', 2310)
('snack', 2295)
('know', 2293)
('differ', 2291)
('say', 2289)
('mani', 2244)
('come', 2239)
('easi', 2202)
('milk', 2193)
('hot', 2182)
('organ', 2178)

('alway', 2078)
('everi', 2055)
('fresh', 2045)
('review', 1996)
('wonder', 1951)
('ship', 1931)
('got', 1925)
('calori', 1902)
('local', 1892)
('stuff', 1880)
('right', 1877)
('qualiti', 1870)
('doe', 1861)
('healthi', 1838)
('oil', 1837)
('regular', 1826)
('befor', 1799)
('natur', 1775)
('ingredi', 1771)
('small', 1752)
('definit', 1743)
('did', 1737)
('cooki', 1734)
('someth', 1717)
('size', 1710)
('sauc', 1705)
('ad', 1688)
('bread', 1680)
('cook', 1671)
('feel', 1637)
('hard', 1637)
('tasti', 1631)
('quick', 1624)
('excel', 1605)
('sure', 1586)
('dark', 1579)
('doesnt', 1578)
('butter', 1572)
('long', 1571)
('chicken', 1561)
('month', 1551)
('happi', 1548)

('help', 1529)
('serv', 1521)
('didnt', 1508)
('strong', 1497)
('actual', 1484)
('big', 1484)
('start', 1483)
('far', 1472)
('contain', 1471)
('quit', 1438)
('problem', 1437)
('howev', 1433)
('thank', 1433)
('rice', 1431)
('old', 1428)
('bake', 1409)
('dri', 1403)
('pretti', 1390)
('bottl', 1386)
('usual', 1375)
('diet', 1373)
('peanut', 1356)
('textur', 1355)
('low', 1352)
('fat', 1348)
('salt', 1347)
('candi', 1332)
('smell', 1331)
('bean', 1318)
('chip', 1291)
('peopl', 1286)
('open', 1244)
('friend', 1239)
('real', 1227)
('groceri', 1221)
('green', 1220)
('fruit', 1212)
('juic', 1212)
('famili', 1207)
('anoth', 1195)
('blend', 1191)
('ice', 1191)

```
('anyth', 1187)
('new', 1174)
('worth', 1169)
('chew', 1165)
('varieti', 1165)
('gluten', 1162)
('roast', 1161)
('protein', 1160)
('expens', 1158)
('meal', 1147)
('arriv', 1142)
('minut', 1142)
('nut', 1139)
('save', 1132)
('sever', 1132)
('thought', 1122)
('morn', 1121)
('ginger', 1109)
('item', 1102)
('light', 1101)
('kid', 1097)
('week', 1087)
('recip', 1081)
('expect', 1080)
('came', 1078)
('half', 1075)
('home', 1068)
('onc', 1056)
('prefer', 1049)
('reason', 1048)
('amaz', 1047)
('receiv', 1039)
('larg', 1028)
('theyr', 1028)
('cereal', 1025)
('absolut', 1023)
```

```
In [44]: def PlotWordCloud(frequency):  
    worcloudPlot = WordCloud(background_color="white", width=1500, height=1000)  
    worcloudPlot.generate_from_frequencies(frequencies=frequency)  
    plt.figure(figsize=(15,10))  
    plt.imshow(worcloudPlot, interpolation="bilinear")  
    plt.axis("off")  
    plt.show()
```

```
In [45]: PlotWordCloud(positive_dict)
```



This is a Word Cloud for all the positive reviews in the corpus.

This Word Cloud plot correponds to the most frequent words in all positive reviews.

How I have plotted this word cloud. Now BoW representation is based on the count of the word in a document. If a word W_i occurs in many document then the sum of its column will be high. Therefore, I have just calculated the sum of all the column and created the dictionary where keys are the features and values are the sum of that column. I feeded this dictionary to the wordcloud and plotted the same. the same procedure is followed for negative reviews as well

```
In [46]: negative_bow_vect = CountVectorizer(stop_words = "english")
negative_bow = negative_bow_vect.fit_transform(negative_reviews["ProcessedText"].values)
```

```
In [47]: negative_bow.shape
```

```
Out[47]: (30000, 21669)
```

```
In [48]: features_negative = negative_bow_vect.get_feature_names()
len(features_negative), type(features_negative)
```

```
Out[48]: (21669, list)
```

```
In [49]: count = []
for i in range(len(features_negative)):
    total = negative_bow.getcol(i).sum() # it will give sum of all the values in 'i'th column
    count.append(total)
```

```
In [50]: negative_dict = dict(zip(features_negative, count))
```

```
In [51]: sortedDict_Negative = sorted(negative_dict.items(), key = lambda negative_dict: negative_dict[1], reverse = True)
```

```
In [52]: for i in range(200):  
         print(sortedDict_Negative[i])
```

```
('tast', 18519)  
('like', 17349)  
('product', 14820)  
('flavor', 10387)  
('just', 9860)  
('tri', 9551)  
('veri', 8848)  
('good', 8065)  
('use', 8041)  
('coffe', 7995)  
('buy', 7165)  
('food', 6949)  
('order', 6571)  
('dont', 6387)  
('tea', 5895)  
('box', 5586)  
('becaus', 5462)  
('amazon', 5284)  
('onli', 5245)  
('time', 5220)  
('make', 5208)  
('eat', 5192)  
('bag', 5151)  
('dog', 5103)  
('realli', 5075)  
('look', 4843)  
('love', 4635)  
('packag', 4621)  
('review', 4408)  
('purchas', 4207)  
('did', 4153)  
('bought', 4046)  
('ani', 3995)  
('bad', 3983)  
('better', 3870)  
('chocol', 3841)  
('disappoint', 3812)  
('want', 3798)
```

('drink', 3771)
('water', 3730)
('think', 3694)
('price', 3554)
('say', 3505)
('know', 3490)
('didnt', 3460)
('ingredi', 3370)
('smell', 3268)
('brand', 3202)
('sugar', 3175)
('great', 3141)
('way', 3139)
('thought', 3125)
('got', 3100)
('ive', 3089)
('littl', 3077)
('someth', 3011)
('store', 2877)
('thing', 2863)
('receiv', 2828)
('befor', 2762)
('money', 2755)
('differ', 2729)
('item', 2724)
('open', 2696)
('mix', 2691)
('cup', 2680)
('pack', 2659)
('day', 2626)
('sweet', 2536)
('doe', 2483)
('year', 2467)
('contain', 2374)
('cooki', 2348)
('stuff', 2335)
('howev', 2334)
('sinc', 2319)
('recommend', 2311)
('ship', 2309)
('away', 2307)
('work', 2276)

('compani', 2257)
('old', 2242)
('doesnt', 2219)
('dri', 2192)
('sure', 2161)
('cat', 2139)
('bottl', 2124)
('come', 2088)
('expect', 2076)
('actual', 2065)
('mani', 2064)
('hard', 2063)
('lot', 2042)
('bar', 2040)
('anoth', 1988)
('return', 1986)
('treat', 1985)
('hope', 1964)
('new', 1955)
('qualiti', 1954)
('read', 1932)
('natur', 1910)
('wast', 1904)
('need', 1903)
('problem', 1903)
('mayb', 1859)
('bit', 1827)
('small', 1827)
('star', 1818)
('peopl', 1764)
('whi', 1763)
('wont', 1755)
('noth', 1746)
('high', 1744)
('organ', 1738)
('list', 1702)
('anyth', 1686)
('textur', 1675)
('month', 1655)
('oil', 1652)
('candi', 1644)
('said', 1632)

('arriv', 1593)
('piec', 1567)
('hot', 1557)
('local', 1557)
('real', 1556)
('ill', 1554)
('free', 1544)
('milk', 1523)
('best', 1521)
('green', 1508)
('feel', 1506)
('enjoy', 1504)
('end', 1481)
('salt', 1467)
('label', 1465)
('ad', 1464)
('stick', 1451)
('coconut', 1433)
('right', 1432)
('case', 1426)
('chang', 1421)
('instead', 1418)
('sauc', 1418)
('bitter', 1403)
('strong', 1399)
('pretti', 1386)
('bean', 1372)
('came', 1370)
('sever', 1351)
('big', 1346)
('kind', 1344)
('chew', 1343)
('worth', 1342)
('regular', 1336)
('half', 1331)
('fresh', 1322)
('gave', 1318)
('start', 1314)
('week', 1308)
('far', 1307)
('everi', 1295)
('size', 1294)

```
('chip', 1289)
('color', 1286)
('wasnt', 1270)
('probabl', 1247)
('terribl', 1246)
('fruit', 1244)
('care', 1238)
('horribl', 1234)
('juic', 1231)
('chicken', 1224)
('date', 1222)
('peanut', 1216)
('unfortun', 1210)
('isnt', 1197)
('guess', 1188)
('cereal', 1187)
('aw', 1181)
('plastic', 1180)
('usual', 1177)
('butter', 1175)
('quit', 1175)
('calori', 1170)
('rice', 1170)
('second', 1164)
('corn', 1154)
('throw', 1153)
('notic', 1145)
('went', 1145)
('wouldnt', 1142)
('save', 1134)
('cost', 1130)
('long', 1127)
('mouth', 1126)
('leav', 1125)
('nice', 1125)
('custom', 1121)
```

```
In [53]: PlotWordCloud(negative_dict)
```



This is a Word Cloud for all the negative reviews in the corpus.

This Word Cloud plot corresponds to the most frequent words in all negative reviews.

Converting Rows to binary BoW

```
In [32]: count_vect = CountVectorizer(binary=True)
```

```
In [33]: Final_Data_BoW = count_vect.fit_transform(Final_Data["ProcessedText"].values)
```

```
In [34]: Final_Data_BoW.shape
```

```
Out[34]: (60000, 29748)
```

```
In [35]: print(type(Final_Data_BoW))  
  
<class 'scipy.sparse.csr.csr_matrix'>
```

Applying Forward chaining cross validation to find best value of alpha

```
In [36]: alphaValue = []  
         for i in range(1, 500, 2):  
             alphaValue.append(i)
```

```

In [37]: # Applying forward chaining cross validation on 80% of data, means 48000 rows. Rest 12000 will be test data.
end1 = 1000
end2 = 2000
FinalScoreOfScores = []

for i in range(47):
    train = Final_Data_Bow[0:end1]
    train_labels = Final_Data_Labels[0:end1]
    test = Final_Data_Bow[end1:end2]
    test_labels = Final_Data_Labels[end1:end2]

    scoreOfScores = []
    for a in alphaValue:
        clf = BernoulliNB(alpha=a, binarize=None, class_prior=None, fit_prior=True)
        clf.fit(train, train_labels)
        score = clf.score(test, test_labels)
        scoreOfScores.append(score)
    FinalScoreOfScores.append(scoreOfScores)

    end1 += 1000
    end2 += 1000

```

```

In [42]: #finding best alpha using brute force
maximum = FinalScoreOfScores[0][0]

for i in range(len(FinalScoreOfScores)):
    Alpha = -1
    for j in range(len(FinalScoreOfScores[i])):
        Alpha += 2
        if maximum < FinalScoreOfScores[i][j]:
            maximum = FinalScoreOfScores[i][j]
            best_alpha = Alpha

print("Maximum Score = "+str(maximum))
print("Value of best alpha = "+str(best_alpha))

```

Maximum Score = 0.865
Value of best alpha = 1

```
In [43]: error = []
maxAlpha = []
splitNumber = []

for i in range(len(FinalScoreOfScores)):
    a = -1
    maximum = -10 #any small negative number can be given
    for j in range(len(FinalScoreOfScores[i])):
        a += 2
        if maximum < FinalScoreOfScores[i][j]:
            maximum = FinalScoreOfScores[i][j]
            max_alpha = a


    error.append(1 - maximum)
    maxAlpha.append(max_alpha)
    splitNumber.append(i+1)
```

```
In [44]: plt.plot(maxAlpha, error)

minError1 = min(error)
minError = np.round(minError1, 2)

plt.xlabel("Alpha")
plt.ylabel("Error")
for xy in zip(maxAlpha, np.round(error,2)):
    plt.annotate(xy,xy)

plt.show()
```



```
In [45]: for xy in zip(maxAlpha, np.round(error,2), splitNumber):  
         print(xy)
```

```
(1, 0.36, 1)  
(1, 0.39, 2)  
(1, 0.31, 3)  
(1, 0.26, 4)  
(1, 0.26, 5)  
(1, 0.26, 6)  
(1, 0.24, 7)  
(1, 0.2, 8)  
(1, 0.23, 9)  
(1, 0.16, 10)  
(1, 0.21, 11)  
(1, 0.17, 12)  
(1, 0.19, 13)  
(1, 0.18, 14)  
(1, 0.18, 15)  
(1, 0.19, 16)  
(1, 0.19, 17)  
(1, 0.19, 18)  
(1, 0.18, 19)  
(1, 0.17, 20)  
(1, 0.19, 21)  
(1, 0.23, 22)  
(1, 0.18, 23)  
(1, 0.18, 24)  
(1, 0.18, 25)  
(1, 0.17, 26)  
(1, 0.15, 27)  
(1, 0.17, 28)  
(1, 0.17, 29)  
(1, 0.17, 30)  
(1, 0.15, 31)  
(1, 0.14, 32)  
(1, 0.16, 33)  
(1, 0.16, 34)  
(1, 0.15, 35)  
(1, 0.14, 36)  
(1, 0.15, 37)  
(1, 0.16, 38)
```


(1, 0.16, 39)
(1, 0.14, 40)
(1, 0.14, 41)
(1, 0.19, 42)
(1, 0.17, 43)
(1, 0.15, 44)
(1, 0.17, 45)
(3, 0.17, 46)
(1, 0.16, 47)

In forward chaining cross validation that has been applied above--in this dataset--contains 47 splits.

Splits are as follows:

train0: 0:1000 test0: 1000:2000

train1: 0:2000 test1: 2000:3000

train2: 0:3000 test2: 3000:4000

train3: 0:4000 test3: 4000:5000

..

..

..

train46: 0:47000 test46: 47000:48000

Now from every split, maximum score is taken and corresponding "alpha" value is taken. So, we will have 47 maximum scores and 47 corresponding alpha values.

Finally, we have plotted them above.

We can also see above that error is decreasing from split1 towards split47

Applying Naive Bayes

```
In [46]: clf = BernoulliNB(alpha=best_alpha, binarize=None, class_prior=None, fit_prior=True)
         clf.fit(Final_Data_Bow[0:48000], Final_Data_Labels[0:48000])

         prediction = clf.predict(Final_Data_Bow[48000:60000])

         score = clf.score(Final_Data_Bow[48000:60000], Final_Data_Labels[48000:60000])
         print(score)

0.8388333333333333
```

Accuracy

```
In [47]: Accuracy = accuracy_score(Final_Data_Labels[48000:60000], prediction) * 100
         print("Accuracy = "+str(Accuracy)+"%")

Accuracy = 83.88333333333333%
```

Precision

```
In [48]: Precision = precision_score(Final_Data_Labels[48000:60000], prediction)
         print("Precision Score = "+str(Precision))

Precision Score = 0.7919807538091419
```

Recall

```
In [49]: Recall = recall_score(Final_Data_Labels[48000:60000], prediction)
         print("Recall Score = "+str(Recall))

Recall Score = 0.8857399103139013
```

F1 Score

```
In [50]: F1_Score = f1_score(Final_Data_Labels[48000:60000], prediction)
print("F1 Score = "+str(F1_Score))
```

F1 Score = 0.8362404741744284

Confusion Matrix

```
In [51]: Confusion_Matrix = confusion_matrix(Final_Data_Labels[48000:60000], prediction)
print("Confusion Matrix \n"+str(Confusion_Matrix))
```

```
Confusion Matrix
[[5128 1297]
 [ 637 4938]]
```

```
In [52]: tn, fp, fn, tp = confusion_matrix(Final_Data_Labels[48000:60000], prediction).ravel()
tn, fp, fn, tp
```

Out[52]: (5128, 1297, 637, 4938)

Getting top 10 Best features

```
In [56]: featuresNames = count_vect.get_feature_names()
len(featuresNames)
```

Out[56]: 29748

```
In [58]: featuresProbabilities = clf.feature_log_prob_
featuresProbabilities.shape
```

Out[58]: (2, 29748)

```
In [59]: featuresProbabilitiesTrans = featuresProbabilities.T
featuresProbabilitiesTrans.shape
```

Out[59]: (29748, 2)

```
In [74]: features_prob_dataFrame = pd.DataFrame(featuresProbabilitiesTrans, index = None, columns = ["Class 0 Probability Scores",
features_prob_dataFrame["Feature Names"] = featuresNames
```

```
In [77]: impFeatures_Class0 = features_prob_dataFrame.sort_values(by = 'Class 0 Probability Scores', axis = 0, ascending = False)
impFeatures_Class0.drop(["Class 1 Probability Scores"], axis = 1, inplace = True)
impFeatures_Class0.head(10)
```

```
Out[77]:
```

	Class 0 Probability Scores	Feature Names
26097	-0.130283	the
963	-0.260555	and
26231	-0.382136	this
17661	-0.582558	not
3558	-0.688957	but
9798	-0.705567	for
28562	-0.740971	was
26085	-0.742930	that
11737	-0.922865	have
25768	-0.952217	tast

Here, above we got all the probability scores of all the features given class label 0. They are nothing but "likelihood" probabilities of all the features given class label 0. Here, in above dataframe values corresponding to column "Class 0" has all the probability scores of features given class 0. The above dataframe has been sorted according to the probability scores of Class 0. On Features column we got important feature names for points belong to class label 0.

```
In [78]: impFeatures_Class1 = features_prob_dataFrame.sort_values(by = 'Class 1 Probability Scores', axis = 0, ascending = False)
impFeatures_Class1.drop(["Class 0 Probability Scores"], axis = 1, inplace = True)
impFeatures_Class1.head(10)
```

```
Out[78]:
```

	Class 1 Probability Scores	Feature Names
26097	-0.196612	the
963	-0.202510	and
26231	-0.450536	this
9798	-0.607925	for
11737	-0.900229	have
29103	-0.910963	with
26085	-0.917089	that
3558	-0.936325	but
1297	-1.108279	are
29536	-1.108651	you

Here, above we got all the probability scores of all the features given class label 1. They are nothing but "likelihood" probabilities of all the features given class label 1. Here, in above dataframe values corresponding to column "Class 1" has all the probability scores of features given class 1. The above dataframe has been sorted according to the probability scores of Class 1. On Features column we got important feature names for points belong to class label 1.

TFIDF

```
In [79]: positiveReviews_5000 = allPositiveReviews[:5000]
```

```
In [80]: positiveReviews_5000.shape
```

```
Out[80]: (5000, 12)
```

```
In [81]: positiveReviews_5000.head()
```

Out[81]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0	0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality Dog Food
2	2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all
3	4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	1	1350777600	Great taffy
4	5	6	B006K2ZZ7K	ADT0SRK1MGOEU	Twoapennything	0	0	1	1342051200	Nice Taffy
5	6	7	B006K2ZZ7K	A1SP2KVKFXXRU1	David C. Sullivan	0	0	1	1340150400	Great! Just as good as the expensive brands!

```
In [82]: negativeReviews_5000 = allNegativeReviews[:5000]
```

```
In [83]: negativeReviews_5000.shape
```

Out[83]: (5000, 12)

```
In [84]: negativeReviews_5000.head()
```

Out[84]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	T	
	1	1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertised	Prod arriv labe
	11	12	13	B0009XLVG0	A327PCT23YH90	LT	1	1	0	1339545600	My Cats Are Not Fans of the New Food	Jun Sal Peanu My c hi be hap eal Felic Pla
	15	16	17	B001GVISJM	A3KLWF6WQ5BNYO	Erica Neathery	0	0	0	1348099200	poor taste	I l eal th and tl are gr
	25	26	27	B001GVISJM	A3RXAU2N8KV45G	lady21	0	1	0	1332633600	Nasty No flavor	watc - cand just re No fla . J pla
	45	47	51	B001EO5QW8	A108P30XVUFKXY	Roberto A	0	7	0	1203379200	Don't like it	T oatm is good. mus so

```
In [85]: frames_10000 = [positiveReviews_5000, negativeReviews_5000]
```

```
In [86]: FinalPositiveNegative = pd.concat(frames_10000)
```

```
In [87]: FinalPositiveNegative.shape
```

```
Out[87]: (10000, 12)
```

```
In [88]: #Sorting FinalDataframe by "Time"  
FinalSortedPositiveNegative_10000 = FinalPositiveNegative.sort_values('Time', axis=0, ascending=True, inplace=False)
```

```
In [89]: FinalSortedPositiveNegativeScore_10000 = FinalSortedPositiveNegative_10000["Score"]
```

```
In [90]: FinalSortedPositiveNegative_10000.shape
```

```
Out[90]: (10000, 12)
```

```
In [91]: FinalSortedPositiveNegativeScore_10000.shape
```

```
Out[91]: (10000,)
```



```
In [92]: FinalSortedPositiveNegative_10000.head()
```

```
Out[92]:
```

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
772	1146	1245	B00002Z754	A29Z5PI9BW2PU3	Robbie	7	7	1	961718400	Great Product
771	1145	1244	B00002Z754	A3B8RCEI0FXFI6	B G Chase	10	10	1	962236800	WOW Make your own 'slickers' !
5822	7427	8111	B0000EIE2Z	A3M174IC0VXOS2	Gail Cooke	3	3	1	1075420800	BEST BLUEBERRIES
2418	3481	3783	B00016UX0K	AF1PV3DIC0XM7	Robert Ashton	1	2	1	1081555200	Classic Condiment
5206	6790	7432	B0001E1IME	A2IKCTD1I73PLW	Adeba	2	8	1	1083456000	amazon monopoly/ripoff



```
In [93]: Final_Data = FinalSortedPositiveNegative_10000
```

```
In [94]: Final_Data.shape
```

```
Out[94]: (10000, 12)
```

```
In [95]: Final_Data_Labels = FinalSortedPositiveNegativeScore_10000
```

```
In [96]: Final_Data_Labels.shape
```

```
Out[96]: (10000,)
```

```
In [72]: positive_reviews = Final_Data[(Final_Data["Score"] == 1)]  
negative_reviews = Final_Data[(Final_Data["Score"] == 0)]
```

```
In [73]: Positive_tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), stop_words = "english")  
Positive_tf_idf = Positive_tf_idf_vect.fit_transform(positive_reviews["ProcessedText"].values)
```

```
In [74]: Positive_tf_idf.shape
```

```
Out[74]: (5000, 114651)
```

```
In [75]: features = Positive_tf_idf_vect.get_feature_names()
```

```
In [76]: idfValues = Positive_tf_idf_vect.idf_
```

```
In [78]: d = dict(zip(features, 9 - idfValues))
```

```
In [79]: sortedDict = sorted(d.items(), key = lambda d: d[1], reverse = True)
```

```
In [80]: for i in range(200):  
         print(sortedDict[i])
```

```
('like', 6.809072442421418)  
('tast', 6.765368008186689)  
('great', 6.763992492151379)  
('good', 6.73466078043391)  
('love', 6.689725684788852)  
('flavor', 6.590032302691801)  
('use', 6.568508292946707)  
('just', 6.50636578331954)  
('veri', 6.503690792870237)  
('product', 6.470097075582087)  
('tri', 6.467323148699362)  
('make', 6.391361607896317)  
('buy', 6.114608605976726)  
('onli', 6.072907876777782)  
('time', 6.060468186302143)  
('realli', 6.049279258384337)  
('best', 5.968767617525185)  
('price', 5.95030555468545)  
('littl', 5.909095286038787)  
('dont', 5.886181026515912)  
('order', 5.886181026515912)  
('amazon', 5.849077276312534)  
('store', 5.833492545295836)  
('eat', 5.817661080079155)  
('coffe', 5.8141086784747875)  
('becaus', 5.806965790962407)  
('recommend', 5.797964830103432)  
('better', 5.790705270090627)  
('mix', 5.726773729244833)  
('ive', 5.650123319469437)  
('ani', 5.631075124498743)  
('high', 5.624644234168453)  
('food', 5.594074168083775)  
('bag', 5.591854411345462)  
('drink', 5.580681110747337)  
('year', 5.569381555493403)  
('delici', 5.544063747509114)  
('want', 5.529979007627374)
```

('work', 5.515693050379898)
('look', 5.501200043077331)
('favorit', 5.474071375689078)
('pack', 5.47156824547096)
('nice', 5.458957737879031)
('enjoy', 5.448753567704788)
('sweet', 5.441031521610878)
('day', 5.433249381168823)
('purchas', 5.428027437187671)
('sugar', 5.3879686766356665)
('say', 5.379760696217836)
('tea', 5.379760696217836)
('perfect', 5.377009662845947)
('snack', 5.377009662845947)
('brand', 5.371484786913977)
('bought', 5.3603426103607354)
('easi', 5.357537559433126)
('way', 5.349074885514393)
('sinc', 5.337678750783524)
('mani', 5.3290456036388205)
('thing', 5.3290456036388205)
('fresh', 5.302689758933458)
('need', 5.302689758933458)
('cup', 5.281699483041622)
('differ', 5.281699483041622)
('bit', 5.278664579346469)
('lot', 5.278664579346469)
('box', 5.269504209947804)
('think', 5.269504209947804)
('free', 5.254047951711112)
('packag', 5.247797931365941)
('add', 5.2094545761682935)
('come', 5.2094545761682935)
('regular', 5.2094545761682935)
('wonder', 5.1963396340904655)
('chip', 5.189717093329972)
('water', 5.189717093329972)
('alway', 5.166186595919777)
('local', 5.1627794375981635)
('everi', 5.155930095752589)
('got', 5.152487751561615)
('right', 5.135096008849747)

('know', 5.131581066742303)
('ship', 5.12451389951921)
('tasti', 5.120961497914842)
('calori', 5.113818610402461)
('review', 5.077318208182936)
('definit', 5.066103137362795)
('hard', 5.058555931727413)
('excel', 5.050951332342193)
('did', 5.039434890280634)
('befor', 5.031682913476317)
('qualiti', 5.027784273060659)
('healthi', 5.000059725045804)
('natur', 5.000059725045804)
('sure', 5.000059725045804)
('stuff', 4.975668271921645)
('famili', 4.971544554737783)
('high recommend', 4.967403762071751)
('someth', 4.963245751923088)
('ad', 4.954877502252571)
('doe', 4.954877502252571)
('chocol', 4.950666969716227)
('hot', 4.950666969716227)
('long', 4.933645282146797)
('happi', 4.929344200247407)
('small', 4.911952457535537)
('thank', 4.911952457535537)
('treat', 4.898707230785517)
('doesnt', 4.894252880436136)
('milk', 4.894252880436136)
('ingredi', 4.876234374933458)
('serv', 4.867101891370185)
('actual', 4.857885236265261)
('feel', 4.857885236265261)
('far', 4.853244856708759)
('howev', 4.843898994290521)
('month', 4.839193103253109)
('problem', 4.839193103253109)
('size', 4.839193103253109)
('didnt', 4.824941080545907)
('quick', 4.820144908282414)
('light', 4.800726822425313)
('usual', 4.800726822425313)

('pancak', 4.795812807622884)
('veri good', 4.795812807622884)
('help', 4.7859117366401716)
('groceri', 4.780924195129133)
('expens', 4.770873859275632)
('varieti', 4.770873859275632)
('start', 4.750464987644424)
('contain', 4.745297017485981)
('old', 4.745297017485981)
('pretti', 4.745297017485981)
('salt', 4.745297017485981)
('quit', 4.740102200608877)
('bake', 4.734880256627726)
('cook', 4.724353843640738)
('friend', 4.724353843640738)
('big', 4.719048791411046)
('gluten', 4.719048791411046)
('sauc', 4.7137154454356835)
('came', 4.7083535022942975)
('worth', 4.697542586190082)
('juic', 4.692092981422517)
('thought', 4.686613515657892)
('dog', 4.664390378873181)
('pleas', 4.664390378873181)
('recip', 4.664390378873181)
('peopl', 4.653090823619248)
('reason', 4.653090823619248)
('open', 4.64739280250461)
('arriv', 4.641662127795625)
('bean', 4.641662127795625)
('kid', 4.641662127795625)
('low', 4.641662127795625)
('abl', 4.635898423078875)
('anyth', 4.635898423078875)
('home', 4.635898423078875)
('case', 4.6301013053945494)
('new', 4.6301013053945494)
('altern', 4.618405265631358)
('groceri store', 4.618405265631358)
('strong', 4.606570807984355)
('expect', 4.600600640997851)
('diet', 4.5945946169376395)

```
('save', 4.588552302481677)
('tast like', 4.588552302481677)
('husband', 4.582473256405295)
('real', 4.582473256405295)
('anoth', 4.576357029387859)
('bad', 4.576357029387859)
('absolut', 4.57020316381348)
('organ', 4.57020316381348)
('tast great', 4.564011193565559)
('item', 4.551511030801327)
('especi', 4.538852633929404)
('smell', 4.532462835830634)
('textur', 4.532462835830634)
('chicken', 4.519559430994725)
('instead', 4.519559430994725)
('sever', 4.519559430994725)
('week', 4.513044749973531)
('oil', 4.506487349427372)
('amaz', 4.499886665396021)
('blend', 4.486553134526556)
('fruit', 4.486553134526556)
('theyr', 4.486553134526556)
('bottl', 4.479819102345211)
('butter', 4.479819102345211)
('kind', 4.479819102345211)
('onc', 4.479819102345211)
('orang', 4.479819102345211)
('wont', 4.459340571001671)
('cat', 4.445451458841004)
('fat', 4.445451458841004)
('prefer', 4.445451458841004)
('ago', 4.438433886182358)
('avail', 4.438433886182358)
('gluten free', 4.438433886182358)
('soda', 4.438433886182358)
('wish', 4.438433886182358)
```

```
In [82]: PlotWordCloud(d)
```



This is a Word Cloud for all the positive reviews in the corpus.

This Word Cloud plot corresponds to the most frequent words in all positive reviews based on IDF values. More is the IDF Value for a word the less frequent is the word in the corpus.

How I have plotted this word cloud. Now formulae for $IDF(D, W_i) = (\ln(N+1 / n_i+1) + 1)$ where 'N' is total number of documents in a corpus and 'ni' is the total number of documents where word 'Wi' occurs. Hence, I got all the idf values from idf_ attribute and I got corresponding features from get_features_names() function. Now since, the highest possible idf value can be 8.88, hence, I subtracted all the idf values from '9' which leads to the highest idf value of the most frequently occurring word. Now I created dictionary where features are the keys and modified idf value are the values and I feeded this to the word cloud and plot the same. The same is done for negative reviews as well.

```
In [83]: Negative_tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), stop_words = "english")  
Negative_tf_idf = Negative_tf_idf_vect.fit_transform(negative_reviews["ProcessedText"].values)
```

```
In [84]: Negative_tf_idf.shape
```

```
Out[84]: (5000, 139161)
```

```
In [85]: features_neg = Negative_tf_idf_vect.get_feature_names()
```

```
In [86]: negIDF = Negative_tf_idf_vect.idf_
```

```
In [87]: NegD = dict(zip(features_neg, 9 - negIDF))
```

```
In [88]: sortedDictNeg = sorted(NegD.items(), key = lambda NegD: NegD[1], reverse = True)
```

```
In [89]: for i in range(200):  
         print(sortedDictNeg[i])
```

```
('tast', 7.0050070599682215)  
('like', 7.004466080782725)  
('product', 6.762615081465285)  
('just', 6.6496447054933165)  
('veri', 6.595748937288184)  
('tri', 6.565994676206391)  
('good', 6.44957396719508)  
('flavor', 6.4170040385096545)  
('use', 6.376263183183732)  
('buy', 6.343270500029383)  
('dont', 6.271578571573267)  
('becaus', 6.170961542527858)  
('onli', 6.1211746177476165)  
('make', 6.113290214223468)  
('time', 6.082477327793933)  
('order', 6.054889371275104)  
('realli', 6.047871798616457)  
('look', 6.009101688151887)  
('love', 5.951857145376868)  
('food', 5.95030555468545)  
('eat', 5.9314962227279535)  
('bought', 5.923553369214017)  
('disappoint', 5.889486814650411)  
('ani', 5.862729365480861)  
('packag', 5.861033012232683)  
('better', 5.8473575854330075)  
('review', 5.8473575854330075)  
('did', 5.835236224900663)  
('think', 5.828243189409692)  
('purchas', 5.822966132308848)  
('amazon', 5.810543612310291)  
('box', 5.794341637734011)  
('want', 5.783392623244341)  
('bad', 5.781556075437039)  
('say', 5.751703112287357)  
('didnt', 5.742188292646019)  
('know', 5.7287135940626595)  
('thought', 5.720931453620604)
```


('bag', 5.713088276159578)
('drink', 5.674969318055968)
('great', 5.66881545248159)
('way', 5.645921632615737)
('price', 5.618171719662835)
('littl', 5.613833318064237)
('got', 5.609476012695281)
('ive', 5.5761765986262315)
('someth', 5.5761765986262315)
('coffe', 5.564825738957542)
('tast like', 5.557952859669781)
('thing', 5.555651362681502)
('water', 5.548714918684844)
('dog', 5.544063747509114)
('store', 5.520477748503234)
('mix', 5.513292088842359)
('money', 5.5036301779306225)
('befor', 5.484021706542246)
('receiv', 5.484021706542246)
('differ', 5.456416440450358)
('open', 5.443612168204369)
('smell', 5.443612168204369)
('away', 5.441031521610878)
('brand', 5.438444198045927)
('doe', 5.4358501628688805)
('ingredi', 5.430641817761742)
('recommend', 5.409532854551507)
('tea', 5.3879686766356665)
('item', 5.382504182163588)
('sugar', 5.382504182163588)
('howev', 5.365929217069375)
('sinc', 5.357537559433126)
('cup', 5.349074885514393)
('old', 5.337678750783524)
('sweet', 5.331931608527956)
('come', 5.314489305864614)
('doesnt', 5.314489305864614)
('year', 5.311552446191303)
('pack', 5.305652724064116)
('work', 5.302689758933458)
('day', 5.29074931856154)
('sure', 5.284725203958159)

('need', 5.266432010910833)
('chocol', 5.263350344373426)
('mani', 5.250927824374869)
('actual', 5.244658211361273)
('hope', 5.222399740760331)
('anoth', 5.219179126060288)
('expect', 5.2094545761682935)
('mayb', 5.193033845955966)
('hard', 5.189717093329972)
('lot', 5.183050401971783)
('ship', 5.166186595919777)
('stuff', 5.159360630849378)
('problem', 5.155930095752589)
('return', 5.152487751561615)
('compani', 5.149033516693528)
('contain', 5.145567308717043)
('wast', 5.135096008849747)
('enjoy', 5.131581066742303)
('anyth', 5.117396431750346)
('bit', 5.106624334768435)
('high', 5.0993779262476675)
('month', 5.095734934969167)
('peopl', 5.088408894877094)
('ill', 5.084725649460797)
('small', 5.077318208182936)
('wont', 5.073593809091953)
('hot', 5.069855486981346)
('noth', 5.069855486981346)
('new', 5.054760860758861)
('star', 5.050951332342193)
('ad', 5.04712723590379)
('read', 5.031682913476317)
('qualiti', 5.019941095599632)
('whi', 5.015996317308616)
('strong', 5.00805976771288)
('treat', 4.996035574746079)
('natur', 4.971544554737783)
('dri', 4.963245751923088)
('textur', 4.963245751923088)
('local', 4.946438633606706)
('right', 4.946438633606706)
('pretti', 4.9379279439387975)

('said', 4.933645282146797)
('kind', 4.929344200247407)
('big', 4.916328832135337)
('gave', 4.916328832135337)
('best', 4.903141827853382)
('real', 4.889778600041215)
('worth', 4.889778600041215)
('chew', 4.885284210453376)
('regular', 4.885284210453376)
('chang', 4.871678558397598)
('feel', 4.871678558397598)
('list', 4.867101891370185)
('half', 4.862504182121556)
('probabl', 4.862504182121556)
('end', 4.839193103253109)
('instead', 4.834464962057163)
('size', 4.834464962057163)
('far', 4.824941080545907)
('came', 4.810482997370677)
('guess', 4.810482997370677)
('arriv', 4.800726822425313)
('free', 4.795812807622884)
('veri disappoint', 4.795812807622884)
('wasnt', 4.795812807622884)
('piec', 4.790874525982302)
('unfortun', 4.780924195129133)
('bitter', 4.770873859275632)
('decid', 4.770873859275632)
('horribl', 4.770873859275632)
('care', 4.765810557319085)
('long', 4.765810557319085)
('everi', 4.7607214878116135)
('sever', 4.7607214878116135)
('start', 4.7607214878116135)
('went', 4.7607214878116135)
('fine', 4.745297017485981)
('usual', 4.740102200608877)
('isnt', 4.734880256627726)
('label', 4.729630900741583)
('case', 4.719048791411046)
('expens', 4.719048791411046)
('aw', 4.7137154454356835)

('notic', 4.7137154454356835)
('throw', 4.7137154454356835)
('bottl', 4.7083535022942975)
('mouth', 4.7083535022942975)
('save', 4.7083535022942975)
('terribl', 4.702962653659421)
('plastic', 4.697542586190082)
('week', 4.697542586190082)
('minut', 4.686613515657892)
('organ', 4.686613515657892)
('wouldnt', 4.686613515657892)
('left', 4.681103859846922)
('complet', 4.675563679471306)
('stick', 4.675563679471306)
('milk', 4.669992634421851)
('let', 4.664390378873181)
('juic', 4.653090823619248)
('second', 4.64739280250461)
('definit', 4.635898423078875)
('took', 4.6301013053945494)
('groceri', 4.61250554350417)
('wast money', 4.61250554350417)
('add', 4.600600640997851)
('low', 4.600600640997851)
('believ', 4.5945946169376395)
('quit', 4.5945946169376395)
('alway', 4.588552302481677)
('onc', 4.588552302481677)
('custom', 4.582473256405295)
('worst', 4.576357029387859)
('ginger', 4.557780643814923)
('nice', 4.557780643814923)
('fact', 4.551511030801327)
('fresh', 4.551511030801327)
('wrong', 4.545201861608064)
('wonder', 4.538852633929404)

```
In [90]: PlotWordCloud(NegD)
```



This is a Word Cloud for all the negative reviews in the corpus.

This Word Cloud plot corresponds to the most frequent words in all negative reviews based on IDF values. More is the IDF Value for a word the less frequent is the word in the corpus.

Converting rows to tfidf

```
In [97]: tfidf_vect = TfidfVectorizer(ngram_range = (1,2))
```

```
In [98]: Final_Data_tfidf = tfidf_vect.fit_transform(Final_Data["ProcessedText"].values)
```

```
In [99]: Final_Data_tfidf.shape
```

```
Out[99]: (10000, 237703)
```

```
In [100]: print(type(Final_Data_tfidf))  
  
<class 'scipy.sparse.csr.csr_matrix'>
```

Applying Forward chaining cross validation to find best value of alpha

```
In [101]: alphaValue = []  
          for i in range(1, 500, 2):  
              alphaValue.append(i)
```



```

In [102]: # Applying forward chaining cross validation on 80% of data, means 8000 rows. Rest 2000 will be test data.
end1 = 1000
end2 = 2000
FinalScoreOfScores = []

for i in range(7):
    train = Final_Data_tfidf[0:end1]
    train_labels = Final_Data_Labels[0:end1]
    test = Final_Data_tfidf[end1:end2]
    test_labels = Final_Data_Labels[end1:end2]

    scoreOfScores = []
    for a in alphaValue:
        clf = MultinomialNB(alpha=a, fit_prior=True, class_prior=None)
        clf.fit(train, train_labels)
        score = clf.score(test, test_labels)
        scoreOfScores.append(score)
    FinalScoreOfScores.append(scoreOfScores)

    end1 += 1000
    end2 += 1000

```

```

In [103]: #finding best alpha using brute force
maximum = FinalScoreOfScores[0][0]

for i in range(len(FinalScoreOfScores)):
    Alpha = -1
    for j in range(len(FinalScoreOfScores[i])):
        Alpha += 2
        if maximum < FinalScoreOfScores[i][j]:
            maximum = FinalScoreOfScores[i][j]
            best_alpha = Alpha

print("Maximum Score = "+str(maximum))
print("Value of best alpha = "+str(best_alpha))

```

Maximum Score = 0.89
Value of best alpha = 1

```
In [104]: error = []
maxAlpha = []
splitNumber = []

for i in range(len(FinalScoreOfScores)):
    a = -1
    maximum = -10 #any small negative number can be given
    for j in range(len(FinalScoreOfScores[i])):
        a += 2
        if maximum < FinalScoreOfScores[i][j]:
            maximum = FinalScoreOfScores[i][j]
            max_alpha = a


    error.append(1 - maximum)
    maxAlpha.append(max_alpha)
    splitNumber.append(i+1)
```

```
In [105]: plt.plot(maxAlpha, error)

minError1 = min(error)
minError = np.round(minError1, 2)

plt.xlabel("Alpha")
plt.ylabel("Error")
for xy in zip(maxAlpha, np.round(error,2)):
    plt.annotate(xy,xy)

plt.show()
```



```
In [106]: for xy in zip(maxAlpha, np.round(error,2), splitNumber):
           print(xy)
```

```
(1, 0.41, 1)
(1, 0.37, 2)
(1, 0.33, 3)
(1, 0.33, 4)
(1, 0.19, 5)
(1, 0.14, 6)
(1, 0.11, 7)
```

In forward chaining cross validation that has been applied above--in this dataset--contains 7 splits.

Splits are as follows:

train0: 0:1000 test0: 1000:2000

train1: 0:2000 test1: 2000:3000

train2: 0:3000 test2: 3000:4000

train3: 0:4000 test3: 4000:5000

..

..

..

train6: 0:7000 test6: 7000:8000

Now from every split, maximum score is taken and corresponding "alpha" value is taken. So, we will have 7 maximum scores and 7 corresponding alpha values.

Finally, we have plotted them above.

We can also see above that error becomes minimum at split 7.

Applying Naive Bayes

```
In [107]: clf = MultinomialNB(alpha=best_alpha, fit_prior=True, class_prior=None)
          clf.fit(Final_Data_tfidf[0:8000], Final_Data_Labels[0:8000])

          prediction = clf.predict(Final_Data_tfidf[8000:10000])

          score = clf.score(Final_Data_tfidf[8000:10000], Final_Data_Labels[8000:10000])
          print(score)
```

0.889

Accuracy

```
In [108]: Accuracy = accuracy_score(Final_Data_Labels[8000:10000], prediction) * 100  
print("Accuracy = "+str(Accuracy)+"%")
```

Accuracy = 88.9%

Precision

```
In [109]: Precision = precision_score(Final_Data_Labels[8000:10000], prediction)  
print("Precision Score = "+str(Precision))
```

Precision Score = 0.8651564185544768

Recall

```
In [110]: Recall = recall_score(Final_Data_Labels[8000:10000], prediction)  
print("Recall Score = "+str(Recall))
```

Recall Score = 0.8921023359288098

F1 Score

```
In [111]: F1_Score = f1_score(Final_Data_Labels[8000:10000], prediction)  
print("F1 Score = "+str(F1_Score))
```

F1 Score = 0.8784227820372398

Confusion Matrix

```
In [112]: Confusion_Matrix = confusion_matrix(Final_Data_Labels[8000:10000], prediction)  
print("Confusion Matrix \n"+str(Confusion_Matrix))
```

Confusion Matrix
[[976 125]
 [97 802]]

```
In [113]: tn, fp, fn, tp = confusion_matrix(Final_Data_Labels[8000:10000], prediction).ravel()
          tn, fp, fn, tp
```

```
Out[113]: (976, 125, 97, 802)
```

Getting top 10 Best features

```
In [115]: featureNames = tfidf_vect.get_feature_names()
          len(featureNames)
```

```
Out[115]: 237703
```

```
In [116]: FeatureProbabilities = clf.feature_log_prob_
          FeatureProbabilities.shape
```

```
Out[116]: (2, 237703)
```

```
In [117]: FeatureProbabilitiesTrans = FeatureProbabilities.T
          FeatureProbabilitiesTrans.shape
```

```
Out[117]: (237703, 2)
```

```
In [119]: features_prob_dataFrame = pd.DataFrame(FeatureProbabilitiesTrans, index = None, columns = ["Class 0 Probability Scores",
          features_prob_dataFrame["Feature Names"] = featureNames
```

```
In [120]: impFeatures_Class0 = features_prob_dataFrame.sort_values(by = 'Class 0 Probability Scores', axis = 0, ascending = False)
impFeatures_Class0.drop(["Class 1 Probability Scores"], axis = 1, inplace = True)
impFeatures_Class0.head(10)
```

```
Out[120]:
```

	Class 0 Probability Scores	Feature Names
197889	-7.130319	the
8036	-7.672842	and
206446	-7.873337	this
133366	-7.986996	not
222275	-8.018136	was
196540	-8.174226	that
76371	-8.212109	for
30555	-8.226640	but
192872	-8.238006	tast
113416	-8.367527	like

Here, above we got all the probability scores of all the features given class label 0. They are nothing but "likelihood" probabilities of all the features given class label 0. Here, in above dataframe values corresponding to column "Class 0" has all the probability scores of features given class 0. The above dataframe has been sorted according to the probability scores of Class 0. On Features column we got important feature names for points belong to class label 0.

```
In [121]: impFeatures_Class1 = features_prob_dataframe.sort_values(by = 'Class 1 Probability Scores', axis = 0, ascending = False)
impFeatures_Class1.drop(["Class 0 Probability Scores"], axis = 1, inplace = True)
impFeatures_Class1.head(10)
```

```
Out[121]:
```

	Class 1 Probability Scores	Feature Names
197889	-7.290579	the
8036	-7.456165	and
206446	-7.951256	this
76371	-8.020146	for
13193	-8.203751	are
235852	-8.263928	you
92550	-8.299243	have
87443	-8.314475	great
196540	-8.343133	that
230790	-8.345903	with

Here, above we got all the probability scores of all the features given class label 1. They are nothing but "likelihood" probabilities of all the features given class label 1. Here, in above dataframe values corresponding to column "Class 1" has all the probability scores of features given class 1. The above dataframe has been sorted according to the probability scores of Class 1. On Features column we got important feature names for points belong to class label 1.